

Redis to the Rescue?

O'Reilly MySQL Conference
2011-04-13



Who?

- Tim Lossen / @tlossen
- Berlin, Germany
- backend developer at wooga





Home

Games

Jobs

About

Press

Support



Diamond Dash


60 seconds gem rush

wooga is now no. 5 worldwide
over 18 m active users every month



Happy Hospital

Cure cute pets from funny diseases



Check out Diamond Dash »



wooga on Facebook



wooga Due to the slow growth of our games on Facebook we have decided that as of today (2011-04-01) we will completely remove our games Monster World, Bubble Island, Diamond Dash, Happy Hospital and Brain Buddies from Facebook and instead focus all of our efforts on developing games for Myspace about an hour ago

wooga Peggy Beschnitt sprach heute bei der Munich Gaming über "Dos and Don'ts bei der Bewerbung und dem Vorstellungsgespräch" für Young Game Talents.
20 hours ago

wooga Diamond Dash, wooga's fifth game, launched two weeks ago and already has more than 1,000,000 users: <http://goo.gl/ntjsw>



Two Weeks, and 1,000,000 Players - wooga's Dash to the Top | wooga

goo.gl

Two numbers that define the wave of explosive success wooga is riding. Diamond Dash, our fifth and newest game, has experienced astonishing growth, capturing

Based in Berlin, wooga is the leading European social games developer.



We are hiring!

Currently, we are searching for:

- [Internship Social Media and PR](#) (m/f)
- [Game Designer - Bubble Island](#) (m/f) (m/f)
- [Software Engineer - Graduate Position](#) (m/f)
- [Internship IT Management](#) (m/f)

Redis Intro

Case 1: Monster World

Case 2: Happy Hospital

Discussion



Redis Intro

Case 1: Monster World

Case 2: Happy Hospital


Discussion



What?


- key-value-store
- in-memory database
- “data structure server”





Data Types

- strings (integers)



Data Types

- strings (integers)
- lists
- hashes



Data Types


- strings (integers)
- lists
- hashes
- sets
- sorted sets





Performance

- same for reads / writes



Performance

- same for reads / writes
- 50 K ops/second
 - regular notebook, EC2 instance




Performance

- same for reads / writes
- 50 K ops/second
 - regular notebook, EC2 instance
- 200 K ops/second
 - intel core i7 X980 (3.33 GHz)



Durability

- snapshots
- append-only log



Other Features

- master-slave replication
- virtual memory
- ...



Redis Intro

Case 1: Monster World

Case 2: Happy Hospital

Discussion



2M people

MonsterWorld

3388

2

+ Add Coins

+ Add WooGoo

9 358

Team wooga

00:30



00:58:17



#57 Sebastian 8 464 ★

#58 Kim 8 433 ★

#59 Juha 8 377 ★

#60 Tim 9 358 ★


#61 Steffen 7 357 ★




Plants

Items








Daily Active Users


Challenge

- traffic growing rapidly



Challenge

- traffic growing rapidly
- bottleneck: write throughput
 - EBS volumes pretty slow




Challenge

- traffic growing rapidly
- bottleneck: write throughput
 - EBS volumes pretty slow
- MySQL already sharded
 - $4 \times 2 = 8$ shards



Idea

- move write-intensive data to Redis



Idea

- move write-intensive data to Redis
- first candidate: inventory
 - integer fields
 - frequently changing



2M people

MonsterWorld

3388

2


+ Add Coins

+ Add WooGoo

9 358

Team wooga

00:30



- #57 Sebastian 464 ★ 8
- #58 Kim 433 ★ 8
- #59 Juha 377 ★ 8
- #60 Tim 358 ★ 9
- #61 Steffen 357 ★ 7
- Invite

Plants Items



Solution

- inventory = Redis hash
 - atomic increment / decrement !



Solution

- inventory = Redis hash
 - atomic increment / decrement !
- on-demand migration of users
 - with batch roll-up



Results

- quick win
 - implemented in 2 weeks
 - 10% less load on MySQL servers




Results

- quick win
 - implemented in 2 weeks
 - 10% less load on MySQL servers
- decision: move over more data



But ...

- “honeymoon soon over”




But ...

- “honeymoon soon over”
- growing memory usage (fragmentation)
 - servers need periodic “refresh”
 - replication dance



Current Status

- hybrid setup
 - 4 MySQL master-slave pairs
 - 4 Redis master-slave pairs



Current Status

- hybrid setup
 - 4 MySQL master-slave pairs
 - 4 Redis master-slave pairs
- evaluating other alternatives
 - Riak, Membase



Redis Intro

Case 1: Monster World

Case 2: Happy Hospital

Discussion



874
 33/14
 19
 7519
 134/200



19 7519

 Tim

19 7456

 Peggy

19 7255

 Stefan

19 7133

 Alena

Invite friend

18 6479

 Anabel


29	10	51	4
17	4	1	4

Challenge

- expected peak load:
 - 16000 concurrent users
 - 4000 requests/second
 - mostly writes




*“Memory is the new Disk,
Disk is the new Tape.”*
— Jim Gray



Idea

- use Redis as main database
 - excellent (write) performance
 - virtual memory for capacity



Idea

- use Redis as main database
 - excellent (write) performance
 - virtual memory for capacity
- no sharding = simple operations



Data Model

- user = single Redis hash
 - each entity stored in hash field (serialized to JSON)



Data Model

- user = single Redis hash
 - each entity stored in hash field (serialized to JSON)
- custom Ruby mapping layer ("Remodel")



```
1 class User < Remodel::Entity
2
3   has_many :pets, :class => Pet
4
5   property :level, :class => Integer, :default => 1
6   property :xp, :class => Integer, :default => 0
7
8 end
9
10
11 class Pet < Remodel::Entity
12
13   property :pet_type, :class => String
14
15 end
```


1220032045	ul	{“level”: 4, “xp”: 241}
	ul_pets	[“p7”, “p8”]
	p7	{“pet_type”: “Cat”}
	p8	{“pet_type”: “Dog”}
1234599660	ul	{“level”: 1, “xp”: 22}
	ul_pets	[“p3”]

1220032045	ul	{“level”: 4, “xp”: 241}
	ul_pets	[“p7”, “p8”]
	p7	{“pet_type”: “Cat”}
	p8	{“pet_type”: “Dog”}
1234599660	ul	{“level”: 1, “xp”: 22}
	ul_pets	[“p3”]

1220032045	ul	{“level”: 4, “xp”: 241}
	ul_pets	[“p7”, “p8”]
	p7	{“pet_type”: “Cat”}
	p8	{“pet_type”: “Dog”}
1234599660	ul	{“level”: 1, “xp”: 22}
	ul_pets	[“p3”]

1220032045	ul	{“level”: 4, “xp”: 241}
	ul_pets	[“p7”, “p8”]
	p7	{“pet_type”: “Cat”}
	p8	{“pet_type”: “Dog”}
1234599660	ul	{“level”: 1, “xp”: 22}
	ul_pets	[“p3”]

1220032045	ul	{“level”: 4, “xp”: 241}
	ul_pets	[“p7”, “p8”]
	p7	{“pet_type”: “Cat”}
	p8	{“pet_type”: “Dog”}
1234599660	ul	{“level”: 1, “xp”: 22}
	ul_pets	[“p3”]

1220032045	ul	{“level”: 4, “xp”: 241}
	ul_pets	[“p7”, “p8”]
	p7	{“pet_type”: “Cat”}
	p8	{“pet_type”: “Dog”}
1234599660	ul	{“level”: 1, “xp”: 22}
	ul_pets	[“p3”]

Virtual Memory

- server: 24 GB RAM, 500 GB disk
 - can only keep “hot data” in RAM




Virtual Memory

- server: 24 GB RAM, 500 GB disk
 - can only keep “hot data” in RAM
- 380 GB swap file
 - 50 mio. pages, 8 KB each




Dec 2010: Crisis

- memory usage growing fast



Dec 2010: Crisis

- memory usage growing fast
- cannot take snapshots any more
 - cannot start new slaves



Dec 2010: Crisis

- memory usage growing fast
- cannot take snapshots any more
 - cannot start new slaves
- random crashes



Analysis

- Redis virtual memory not compatible with:
 - persistence
 - replication



Analysis


- Redis virtual memory not compatible with:
 - persistence
 - replication
- need to implement our own!





Workaround


- “dumper” process
 - tracks active users
 - every 10 minutes, writes them into YAML file














Workaround

- in case of Redis crash
 - start with empty database
 - restore users on demand from YAML files




Real Solution

- Redis “diskstore”
 - keeps all data on disk
 - swaps data into memory as needed



Real Solution


- Redis “diskstore”
 - keeps all data on disk
 - swaps data into memory as needed
- under development (expected Q2)



uptime_in_days:102
total_commands_processed:54428739517
vm_stats_swappin_count:73139647
vm_stats_swappout_count:77343129
db2:keys=4296618,expires=0

Results

- average response time: **10 ms**



Results

- average response time: **10 ms**
- peak traffic:
 - 1500 requests/second
 - 15000 Redis ops/second



Current Status

- very happy with setup
 - simple, robust, fast
 - easy to operate
- still lots of spare capacity



Redis Intro

Case 1: Monster World

Case 2: Happy Hospital

Discussion



Advantages

- order-of-magnitude performance improvement
 - removes main bottleneck
 - enables simple architecture



Disadvantages

- main challenge: durability
 - diskstore very promising




Disadvantages

- main challenge: durability
 - diskstore very promising
- no ad-hoc queries
 - think hard about data model
 - hybrid approach?



Conclusion

- Ruby + Redis = killer combo



Q & A



redis.io

github.com/tlossen/remodel

wooga.com/jobs

